

# DIY Floating Point Addition

This program reads two 32 bit floating point numbers from standard input and adds them together using C's integer operations.

The approach I took to completeing this program was to reuse parts of the last program, A1. I reused the int to binary output function, and by doing so I just had to make sure I outputted the correct number of spaces / newlines. I then added in a chunk of code from the example program, a struct with 32 unsigned bits, which allows me to access specific parts of the mantissa and exponent. I then typedef'd the union so that I could load the two floating point numbers into different unions, and I created the following function:

```
float add_floating_point(float_32 first_int, float_32 second_int);
```

Inside this function I use a couple of if statements to detect a few cases of floating point addition. The exponents are either the same, or different, and if they are different one is larger then the other. So I adjust the exponents so that they are the same, and to do this I increment the smaller exponent while right shifting the smaller exponent's mantissa. This gets the two exponents to be the same, and I deal with certain cases like the hidden bit, shifting too far right, and rounding being off by 1 with additional if statements. I finally have one other case to detect infinite numbers and if I find this, I set the exponent to 255 (1111 1111) and the mantissa to 0 (000 0000 0000 0000 0000 0000).

All of my test cases pass, as well as the provided test cases. I verified this by running them manually and with the Makefile that I created. I compared the output of my program (emulated portion) to what the hardware outputted. The hardware is basically just printing out the addition of two floating point number's using C's standard library.

I learned a fair amount about floating point numbers. For one, they are a pain to deal with manually. There are so many random cases to deal with, and a lot of test cases were required to make sure the program was outputting correctly.

I finally feel that this program deserves a .85, as all of my test cases pass successfully, and the output of my program matches the hardware output on my machine.